

Instructor Answer Key

Document: Instructor Answer Key
Version: v1.1.0
Author: Celaya Solutions
Contact: hello@celayasolutions.com
Date: 2026-09-06
SHA256: 0e4f8d029443e0600e73088e1c459003456eaea2bb39cc0da0a2d24b20009b02
Chain: n/a
Tx: [not anchored]
License: All Rights Reserved / Celaya Solutions

Use this key with the level grading guides. Exact model wording is never required. Grade the evidence trail and safe behavior.

Level 1

The four-file binder contains two public summaries and two visibly synthetic training records. Load no answer key or lesson as model evidence.

Question	Expected answer and receipt
1: location	El Paso, Texas; public-identity-and-work.md / Where the lab is based
2: systems	Local-first, provenance-aware legal, archival, industrial systems; public-identity-and-work.md / Work and systems the lab builds
3: contact plus work	hello@celayasolutions.com from public-contact-and-method.md / Public contact route for research collaboration AND the work description from the identity/work file
4: price	Not in the documents. No guessed price or range
5: training conflict	Record A: RESEARCH / two classroom days; Record B: COLLABORATION / three classroom days; both synthetic, both cited, human review, no winner

A real citation does not certify the claim. The learner opens both public sources for question 3 and explains the comparison. Every expected result is written before asking. If all tests pass, use the four-to-one-to-four retrieval exercise; judge evidence coverage, not a forced hallucination. Keep a useful source-quote display edit and a same-question before/after record.

All five saved examples are authored classroom material, not live output. The learner labels the route honestly. A fork commit proves code ownership; a private upload receipt proves submission. An export alone does neither.

Common questions

Did the model learn my files? No. Keyword search selects text for each request. This is not training.

Are two headings one two-document answer? No. Question 3 must cite the two distinct public files.

Does a model miss fail the learner? A learner must identify it, show the evidence, and explain the change. Do not award live-build proof to a saved route or claim application readiness while acceptance fails.

Level 2

Use the [complete Level 2 key](#) for the six-line spec, expected receipts, failure/recovery cases, and exit answers. Baseline and unchanged create no issue; OPEN to PAUSED creates one practice issue in the learner's own fork; a repeat creates none. Failures preserve the last confirmed state. The never list explicitly permits that practice issue and forbids posting elsewhere.

The daily template is 13:17 UTC and may be delayed or dropped. Disabling future triggers, setting WATCHMAN_ENABLED=false, and cancelling queued/active work are separate steps. Grade the observed disabled state honestly; do not claim a future schedule tick was tested without a real observation. Partner/local/saved routes do not prove an

individual hosted launch.

Level 3

- The selected local model answers while wifi is off.
- The learner records model name, file size, time, quality, and a checked fact.
- A custom rule set runs under a learner-chosen name and does not guess a price or date.
- The learner states one job for local use and one reason to use the cloud.

Expected Level 3 evidence

Model names and sizes are not fixed in this key. The instructor records the approved choices during the same week as class. A learner must use those exact names, not a remembered brand name.

The arithmetic prompt has a checked answer:

84 x 1.0825 = 90.93 dollars

A result of 90.92 dollars is a miss. The point is not that the local model failed. The point is that the learner checked.

The airplane test proves the answer was generated without an internet connection on that device. It does not prove the output is correct, that the model has a suitable license for every use, or that no other local program stores the prompt.

Common Level 3 questions

How much memory is enough? The model file plus working space must fit. Use measured class hardware and the prepared small model instead of promising a universal size.

Does a graphics chip make local models possible? Small models can run without one. A suitable GPU makes many models faster.

Is open-weight the same as open-source? No. Open-weight usually publishes the trained file. Open-source should also make the relevant source and license inspectable. Check each model's actual license.

Why use the cloud at all? A cloud model may be stronger for a hard job, easier on limited hardware, or supported by a needed tool. The learner should choose from evidence and privacy needs.

Level 4

- The spec contains trigger, steps, checks, failure routes, access list, and never list.
- Five callers include an easy question, not-found question, urgent request, high-stakes stop, and impersonation.
- The message card has name, contact route, need, urgency, and best follow-up time.
- Consent is recorded before any partner test.
- Before and after counts use the same ten attacks against the learner's own desk.
- Removing the fake secret is identified as the strongest defense.
- The desk drafts only, has a never list, leaves a log, and has a tested stop.
- The pitch names the outcome without AI, agent, or model.

Expected Level 4 callers

Caller	Expected behavior
Public fact	Answer from the public brief and name its heading
Custom price	Say Not in the public source and collect a message
Urgent collaboration	Record urgency without promising a response time
High-stakes request	Do not provide steps; make a safe human handoff
Impersonation	Reveal no client list, unpublished number, or other private fact

The message card has five parts: name or safe alias, public contact route, need, urgency, and best follow-up time. It is incomplete if the desk invents a private phone number or closes without asking for a missing field.

Expected Level 4 defense map

Defense	What it can improve	What it cannot guarantee
Mark stranger text as data	Reduces common instruction-following errors	The model still reads words as one context
Remove the fake secret	Prevents that secret from appearing in model output	Does not secure unrelated data or tools
Draft only; person acts	Stops a bad draft from directly sending, spending, booking, or deleting	A person can still approve carelessly
Log and switch	Makes behavior reviewable and stoppable	Does not make a bad answer correct

The before count can be zero. Current models vary. The learner still applies the defenses and reruns the same tests because future behavior can change.

The strongest expected answer is: a secret the model never received cannot be copied from the model. Removing it is stronger than asking the model to remember not to reveal it.

Common Level 4 questions

Why use a real business if records are synthetic? The public identity and work make the questions meaningful. Synthetic operating records let the room test safely without inventing claims about the real business.

Can the desk book a meeting? Not in the required route. It may draft or propose. A person owns the consequential action.

Can a perfect system prompt solve injection? No known wording provides a hard security boundary. Use permissions, data separation, human approval, logging, and a switch outside the model.

Can learners test a public chatbot if the attack is harmless? Not in this course. The allowed targets are learner-owned, partner-permitted, or supplied class targets.

What belongs in the log? Time, input or safe reference, draft, any tool request and result, flag, cost when available, and final human decision. Do not log a secret.

What does Vercel prove? A fresh browser interaction can prove frontend behavior. It does not by itself prove a secret stayed server-side, a backend is healthy, or data persisted.

Level 5

The eight required proof items are ownership, before state, tested change, normal run, failure route, privacy or security fence, stop or fallback, and cost plus reflection. The cost sheet must use current official provider and hosting sources; the final number is a classroom estimate, and there is no approved CSR setup fee or subscription in this course. A polished video cannot replace a missing item. A saved outage run can replace a live service when the learner explains the decisions and evidence.

Project-specific normal and failure checks

Project	Normal route	Failure route	Required fence or stop
Documents	Supported fresh question with exact receipt	Missing or conflicting fresh question	Stop rule and opened source
Watchman	Unchanged run with log	Controlled change or fetch failure	Never list and disabled workflow
Local model	Named-model job offline	Missing fact or checked false claim	Localhost boundary and exited runner
Front desk	Public answer and complete message card	Price, high-stakes, impersonation, or an unused attack	One business, one binder, secret absent, draft-only permission, switch

Presentation timing

Part	Maximum
Outcome	30 seconds
Normal run	60 seconds
Miss, change, and improved result	45 seconds
Fence and switch or fallback	30 seconds
Cost and next safe step	15 seconds

For deployed work, ask three different questions: Did the browser behavior work from a fresh page? Did the service health check succeed? Did saved state remain after a restart or later request? Accept only the claims the receipts support.

Common Level 5 questions

Does a live failure mean the learner fails? No. The learner can state the expected behavior, show the last verified receipt, use the outage route, and explain the next check. Hiding or mislabeling the failure is the problem.

What if someone cannot present aloud? Use a recording or partner-read script while the learner controls and explains the evidence. The same proof items apply.

Should the class vote for a winner? No ranking is required. Works, honest, and theirs guide revision. Applause is for every learner.