

Instructor Answer Key

Document: Instructor Answer Key
Version: v1.1.0
Author: Celaya Solutions
Contact: hello@celayasolutions.com
Date: 2026-09-06
SHA256: 29c3bc72285fc3a3e3a29c0e73ada17ff672e00993e493ab3469399e731d1c87
Chain: n/a
Tx: [not anchored]
License: All Rights Reserved / Celaya Solutions

Use this key with the level grading guides. Exact model wording is never required. Grade the evidence trail and safe behavior.

Level 1

The four-file binder contains two public summaries and two visibly synthetic training records. Load no answer key or lesson as model evidence.

Question	Expected answer and receipt
1: location	El Paso, Texas; public-identity-and-work.md / Where the lab is based
2: systems	Local-first, provenance-aware legal, archival, industrial systems; public-identity-and-work.md / Work and systems the lab builds
3: contact plus work	hello@celayasolutions.com from public-contact-and-method.md / Public contact route for research collaboration AND the work description from the identity/work file
4: price	Not in the documents. No guessed price or range
5: training conflict	Record A: RESEARCH / two classroom days; Record B: COLLABORATION / three classroom days; both synthetic, both cited, human review, no winner

A real citation does not certify the claim. The learner opens both public sources for question 3 and explains the comparison. Every expected result is written before asking. If all tests pass, use the four-to-one-to-four retrieval exercise; judge evidence coverage, not a forced hallucination. Keep a useful source-quote display edit and a same-question before/after record.

All five saved examples are authored classroom material, not live output. The learner labels the route honestly. A fork commit proves code ownership; a private upload receipt proves submission. An export alone does neither.

Common questions

Did the model learn my files? No. Keyword search selects text for each request. This is not training.

Are two headings one two-document answer? No. Question 3 must cite the two distinct public files.

Does a model miss fail the learner? A learner must identify it, show the evidence, and explain the change. Do not award live-build proof to a saved route or claim application readiness while acceptance fails.

Level 2

Use the [complete Level 2 key](#) for the six-line spec, expected receipts, failure/recovery cases, and exit answers. Baseline and unchanged create no issue; OPEN to PAUSED creates one practice issue in the learner's own fork; a repeat creates none. Failures preserve the last confirmed state. The never list explicitly permits that practice issue and forbids posting elsewhere.

The daily template is 13:17 UTC and may be delayed or dropped. Disabling future triggers, setting WATCHMAN_ENABLED=false, and cancelling queued/active work are separate steps. Grade the observed disabled state honestly; do not claim a future schedule tick was tested without a real observation. Partner/local/saved routes do not prove an

individual hosted launch.

Level 3: Nothing Leaves the Building

Use the [Level 3 answer key](#). Judge observed or labeled fallback evidence, six fair comparison rows, an independent source and arithmetic check, and both named-helper tests before/after one writing-rule edit. Total completion time is not time to first word. Private proof stays out of Git.

Level 4

Use the [Level 4 answer key](#). Require a six-part spec, five caller verdicts, five-field message card, one citation-rule edit and C1/C2 retests, matched A1-A10 before/after results, four control locations, log and pause proof, and outcome pitch. Keep failed/truncated rows separate from Held. Removing bait from a payload prevents copying it from that context, not guessing or unrelated failures. Saved and mixed evidence stays labeled; an upload needs reopening.

Level 5

The eight required proof items are ownership, before state, tested change, normal run, failure route, privacy or security fence, stop or fallback, and cost plus reflection. The cost sheet must use current official provider and hosting sources; the final number is a classroom estimate, and there is no approved CSR setup fee or subscription in this course. A polished video cannot replace a missing item. A saved outage run can replace a live service when the learner explains the decisions and evidence.

Project-specific normal and failure checks

Project	Normal route	Failure route	Required fence or stop
Documents	Supported fresh question with exact receipt	Missing or conflicting fresh question	Stop rule and opened source
Watchman	Unchanged run with log	Controlled change or fetch failure	Never list and disabled workflow
Local model	Named-model job offline	Missing fact or checked false claim	Localhost boundary and exited runner
Front desk	Public answer and complete message card	Price, high-stakes, impersonation, or an unused attack	One business, one binder, secret absent, draft-only permission, switch

Presentation timing

Part	Maximum
Outcome	30 seconds
Normal run	60 seconds
Miss, change, and improved result	45 seconds
Fence and switch or fallback	30 seconds
Cost and next safe step	15 seconds

For deployed work, ask three different questions: Did the browser behavior work from a fresh page? Did the service health check succeed? Did saved state remain after a restart or later request? Accept only the claims the receipts support.

Common Level 5 questions

Does a live failure mean the learner fails? No. The learner can state the expected behavior, show the last verified receipt, use the outage route, and explain the next check. Hiding or mislabeling the failure is the problem.

What if someone cannot present aloud? Use a recording or partner-read script while the learner controls and explains the evidence. The same proof items apply.

Should the class vote for a winner? No ranking is required. Works, honest, and theirs guide revision. Applause is for every learner.