

Instructor Run Sheet - Project Lab Level 1: Your Documents Answer Back

Document: Instructor Run Sheet - Project Lab Level 1: Your Documents Answer Back
Version: v1.0.0
Author: Celaya Solutions
Contact: hello@celayasolutions.com
Date: 2026-09-06
SHA256: 9a40b06ded623019b8c9cb573ff104bc85f42aa1e667655da34e2fe514d33f18
Chain: n/a
Tx: [not anchored]
License: All Rights Reserved / Celaya Solutions

Outcome

Build and test a custom document helper that searches supplied public and synthetic files, shows exact source receipts, stops on missing facts, and flags conflicts for human review.

Alignment

Target	Practice	Evidence
Own and run a custom app	Fork, locked install, local launch	Fork/commit and provider record
Separate source classes	Classify four files	Two public and two synthetic labels
Evaluate grounded answers	Five predictions and runs	Expected, actual, passages, verdict
Verify a two-document answer	Open both public sources	Two receipts and written comparison
Improve deliberately	One setting change and same-question rerun	Before, change, after
Complete a handoff	Private Markdown upload	Correct file and submission time

Teaching stance

The beginner route is the main route. Teach the binder-and-clerk picture before naming RAG. Do not turn this into a Python syntax lecture. Ask what the learner expects, point to one next step, then return the controls. Grade decisions and evidence, not model speed or subscription tier.

Prepare before learners arrive

1. Follow [the release checklist](#) and run the actual preflight on the cohort's Windows and Mac machines. Record OS, RAM, chip, dependency versions, model names, and results. Do not treat simulated CI as proof of a student's installer experience.
2. Complete the five-question pilot for Ollama, Claude API, and OpenAI API. Save real answers with provider/model/time labels. Recheck the public source at <https://www.celayasolutions.com/> and keep the two public files and answer key aligned.
3. Open the student guide, 20-slide deck, rules card, manual edit card, and complete five-question saved packet. Download the PDFs before class; print large-text worksheets when needed.
4. In the private course platform, use a dedicated test student. Follow the actual sign-in > course > Level 1 navigation, upload a harmless test Markdown file, and reopen its receipt. Record the observed labels/screens in the private release evidence. Never publish access codes or student identity.
5. Confirm Level 1's submission window is open and enrollment invitations are delivered. Provide missing join codes privately. Have a setup clinic for failed preflights.

6. Show a working binder, a missing price, a conflict, and the four-to-one-to-four retrieval exercise. Then show the source-quote improvement that remains in the student's commit.

Safety boundary

Only public and synthetic class files may enter a cloud request. No client records, contracts, private contact, messages, credentials, or unpublished work. Keys and learner data stay in .zta/ and outside public commits. The local app binds to 127.0.0.1; it is not a production public service. Synthetic records are never CSR promises.

90-minute schedule

Block	Minutes
Arrival and live readiness	5
Binder and clerk demonstration	5
Public and synthetic source sort	5
Load binder and build index	8
Write expected results	5
Run five questions and judge answers	15
Open two public source receipts	7
Choose miss or controlled exercise	5
Guided edit and same-question rerun	12
Repeat all five checks	8
Test, review, commit, and push	5
Export and confirm private upload	7
Stop and exit ticket	3
Total	90

Facilitation plan

Opening: show the completed proof first. Identify the coding assistant and answer engine as separate choices. Check that learners already passed preflight; use the saved route after two failed attempts or five minutes.

Demo: trace one question through keyword search, selected passages, model answer, and opened source. Ask the missing-price and synthetic-conflict questions. Point out that the app validates receipts but cannot guarantee every claim is true.

Practice: inspect all five expectations before the first Run. At question 3, require both public files. At question 5, count both labels and targets. Ask the learner to explain why a correct refusal can be a success.

Revision: a real miss is useful. If none appears, use the controlled retrieval limit. Never ask a learner to manufacture an unsupported model answer. Restore four passages, enable source quotes, rerun, and keep the small code change.

Close: verify a public commit and private upload receipt separately. Explain that exporting is not submitting and closing the browser is not stopping the server.

Passing proof

Four classified documents; five predicted and judged tests; two public source checks; one same-question improvement; a fork commit; and a private evidence upload receipt.

Grading guide

Check	Meets	Return for revision
Source classes	All four correct	Synthetic material treated as real
Predictions and tests	Five expected results before tests and five honest verdicts	Missing prediction or mislabeled outcome
Source checks	Both public documents support question 3	Vague citation or unsupported claim accepted
Missing and conflict understanding	No guessed price accepted; both synthetic records identified	Invented price or silent conflict winner accepted
Improvement	Same question, deliberate edit, recorded effect	Different question or hidden failed rerun
Ownership and handoff	Fork commit, private evidence file, receipt	Export alone called submission

A live model miss does not erase good analysis. A complete evaluated packet may accurately report a remaining miss. Classroom-ready APPLICATION acceptance remains all five expected behaviors on every supported engine. Saved routes receive evaluation credit and are labeled saved; they do not count as live build verification.

Access and support

Own Windows/Mac computer is the live build route. Phone/paper learners can evaluate the saved packet and complete their own decisions. Give code-experienced learners the retrieval extension after core proof. Read the English safety line and Spanish anchors aloud. Quiet written sharing is welcome.

Fallbacks

[Saved examples](#) cover all five questions, including the two-document answer. The app imports these only through an explicitly labeled control after predictions are filled. The printed packet includes before/repair examples. Mark live-only fields Not applicable: saved run. Never substitute a prepared answer without labeling it.

After class

Record preparation completion, selected routes, completed proofs, returned proofs, observed app misses, setup/answer times, fallback use, and upload blocks. Keep student data in the private course system. Update release verification with observed evidence. Brief maintenance notes belong in AGENTS.md; do not claim an unrun pilot passed.